

Assignment 3

Kangrui Liu

2023-10-08

You may work in pairs or individually for this assignment. Make sure you join a group in Canvas if you are working in pairs. Turn in this assignment as an HTML or PDF file to ELMS. Make sure to include the R Markdown or Quarto file that was used to generate it. Include the GitHub link for the repository containing these files.

- https://github.com/krlui67/Assignment_SURV727/tree/main/a3

```
library(xml2)
library(rvest)
library(tidyverse)
library(jsonlite)
library(robotstxt)
library(RSocrata)
library(dplyr)
```

Web Scraping

In this assignment, your task is to scrape some information from Wikipedia. We start with the following page about Grand Boulevard, a Chicago Community Area.

https://en.wikipedia.org/wiki/Grand_Boulevard,_Chicago

The ultimate goal is to gather the table “Historical population” and convert it to a `data.frame`.

As a first step, read in the html page as an R object. Extract the tables from this object (using the `rvest` package) and save the result as a new object. Follow the instructions if there is an error. Use `str()` on this new object -- it should be a list. Try to find the position of the “Historical population” in this list since we need it in the next step.

```
link <- "https://en.wikipedia.org/wiki/Grand_Boulevard,_Chicago"
paths_allowed(link)
```

```
## en.wikipedia.org
```

```
## [1] TRUE
```

```
url <- read_html(link)
```

```

# find other links
other_places <- html_element(url, xpath = '//*[@contains(concat( " ", @class, " " ), concat( " ", "navbo
# get links of other places
place_links <- other_places %>% html_nodes("a") %>% html_attr("href") %>% data.frame()
for (i in 1:dim(place_links)[1]) {
  place_links[i,1] <- paste0("https://en.wikipedia.org",place_links[i,1], collapse = NULL)
}
names(place_links)[1] <- "link"
# get names of places
place_names <- other_places %>% html_nodes("a") %>% html_attr("title") %>% data.frame()
names(place_names)[1] <- "name"
place_names <- place_names %>% separate(name,c("place","city"),",")
# combine names and links
links <- cbind(place_links,place_names)
rm(place_links)
rm(place_names)
# combine all links
links[dim(links)[1]+1,] <- c("https://en.wikipedia.org/wiki/Grand_Boulevard,_Chicago","Grand Boulevard"

trimws(links$city)

```

```

## [1] "Chicago"          "Chicago"          "Chicago"
## [4] "Chicago"          "Chicago"          "Chicago (neighborhood)"
## [7] "Chicago"          "Chicago"

```

```
head(links)
```

```

##                               link
## 1 https://en.wikipedia.org/wiki/Armour_Square,_Chicago
## 2 https://en.wikipedia.org/wiki/Douglas,_Chicago
## 3 https://en.wikipedia.org/wiki/Oakland,_Chicago
## 4 https://en.wikipedia.org/wiki/Fuller_Park,_Chicago
## 5 https://en.wikipedia.org/wiki/Kenwood,_Chicago
## 6 https://en.wikipedia.org/wiki/Washington_Park,_Chicago_(neighborhood)
##      place      city
## 1 Armour Square    Chicago
## 2 Douglas          Chicago
## 3 Oakland          Chicago
## 4 Fuller Park      Chicago
## 5 Kenwood          Chicago
## 6 Washington Park Chicago (neighborhood)

```

Extract the “Historical population” table from the list and save it as another object. You can use subsetting via `[[...]]` to extract pieces from a list. Print the result.

You will see that the table needs some additional formatting. We only want rows and columns with actual values (I called the table object `pop`).

```

# scrape data
tables <- data.frame()
for (i in 1:dim(links)[1]){
  url_temp <- read_html(links$link[i])
  html_temp <- html_element(url_temp, xpath = '//*[@contains(concat( " ", @class, " " ), concat( " ", "u

```

```

if(length(html_temp)==0){
  html_temp <- html_element(url_temp, xpath = '//*[contains(concat( " ", @class, " "), concat( " ", 
})
table_temp <- html_table(html_temp)
table_temp$City <- links$city[i]
table_temp$Place <- links$place[i]
table_temp <- table_temp[,-3]
table_temp <- table_temp[-dim(table_temp)[1],]
tables <- rbind(tables,table_temp)
}
# change column name
names(tables)[3] <- "Changes"
names(tables)[2] <- "Pop"
# remove temps
rm(url_temp)
rm(table_temp)
rm(html_temp)

head(tables)

```

```

## # A tibble: 6 x 5
##   Census Pop    Changes City      Place
##   <chr>   <chr>   <chr>   <chr>   <chr>
## 1 1930    21,450 -        " Chicago" Armour Square
## 2 1940    18,472 -13.9%  " Chicago" Armour Square
## 3 1950    23,294 26.1%  " Chicago" Armour Square
## 4 1960    15,783 -32.2%  " Chicago" Armour Square
## 5 1970    13,063 -17.2%  " Chicago" Armour Square
## 6 1980    12,475 -4.5%   " Chicago" Armour Square

```

```

# clean data, turn char to number
tables$Pop <- gsub(",",'',tables$Pop)
for (i in 1:dim(tables)[1]) {
  temp <- tables$Changes[i]
  temp <- gsub("%","",temp)
  if(temp == "-"){
    temp = as.numeric(0.00)
  }else if(str_detect(temp,"-")== TRUE){
    temp <- gsub("-",'',temp)
    temp <- as.numeric(temp)
    temp = -1*temp
    temp <- temp/100
  }else{
    temp <- as.numeric(temp)
    temp <- temp/100
  }
  tables$Changes[i] <- temp
}
rm(temp)

tables$Changes <- as.numeric(tables$Changes)
tables$Pop <- as.numeric(tables$Pop)
head(tables)

```

```
## # A tibble: 6 x 5
##   Census    Pop Changes City      Place
##   <chr>   <dbl>   <dbl> <chr>   <chr>
## 1 1930    21450     0    " Chicago" Armour Square
## 2 1940    18472  -0.139 " Chicago" Armour Square
## 3 1950    23294   0.261 " Chicago" Armour Square
## 4 1960    15783  -0.322 " Chicago" Armour Square
## 5 1970    13063  -0.172 " Chicago" Armour Square
## 6 1980    12475  -0.045 " Chicago" Armour Square
```

Expanding to More Pages

That’s it for this page. However, we may want to repeat this process for other community areas. The Wikipedia page https://en.wikipedia.org/wiki/Grand_Boulevard,_Chicago has a section on “Places adjacent to Grand Boulevard, Chicago” at the bottom. Can you find the corresponding table in the list of tables that you created earlier? Extract this table as a new object.

Then, grab the community areas east of Grand Boulevard and save them as a character vector. Print the result.

We want to use this list to create a loop that extracts the population tables from the Wikipedia pages of these places. To make this work and build valid urls, we need to replace empty spaces in the character vector with underscores. This can be done with `gsub()`, or by hand. The resulting vector should look like this: “Oakland,_Chicago” “Kenwood,_Chicago” “Hyde_Park,_Chicago”

To prepare the loop, we also want to copy our `pop` table and rename it as `pops`. In the loop, we append this table by adding columns from the other community areas.

```
# I actually do this part in my second R block.
```

Build a small loop to test whether you can build valid urls using the vector of places and pasting each element of it after <https://en.wikipedia.org/wiki/> in a for loop. Calling `url` shows the last url of this loop, which should be https://en.wikipedia.org/wiki/Hyde_Park,_Chicago.

Finally, extend the loop and add the code that is needed to grab the population tables from each page. Add columns to the original table `pops` using `cbind()`.

Scraping and Analyzing Text Data

Suppose we wanted to take the actual text from the Wikipedia pages instead of just the information in the table. Our goal in this section is to extract the text from the body of the pages, then do some basic text cleaning and analysis.

First, scrape just the text without any of the information in the margins or headers. For example, for “Grand Boulevard”, the text should start with, “**Grand Boulevard** on the South Side of Chicago, Illinois, is one of the ...”. Make sure all of the text is in one block by using something like the code below (I called my object `description`).

```
# get description of each place
# description <- description %>% paste(collapse = ' ')
for (i in 1:dim(links)[1]){
  url_temp <- read_html(links$link[i])
  html_temp <- html_element(url_temp, xpath = '//p[(((count(preceding-sibling::*) + 1) = 6) and parent:
  if(length(html_temp) == 0){
```

```

    html_temp <- html_element(url_temp, xpath = '//p[(((count(preceding-sibling::*) + 1) = 7) and parent::p[1])]')
  }
  if(length(html_temp) == 0){
    html_temp <- html_element(url_temp, xpath = '//p[(((count(preceding-sibling::*) + 1) = 8) and parent::p[1])]')
  }
  if(length(html_temp) == 0){
    html_temp <- html_element(url_temp, xpath = '//p[(((count(preceding-sibling::*) + 1) = 9) and parent::p[1])]')
  }
  desc_temp <- html_text(html_temp)
  desc_temp <- desc_temp %>% paste(collapse = ' ')
  links$desc[i] <- desc_temp
}
# remove temps
rm(url_temp)
rm(desc_temp)
rm(html_temp)

trimws(links$desc)

```

```

## [1] "Armour Square is a Chicago neighborhood on the city's South Side, as well as a larger, official neighborhood."
## [2] "Douglas, on the South Side of Chicago, Illinois, is one of Chicago's 77 community areas. The neighborhood is located on the South Side of Chicago, Illinois, and is one of the city's 77 community areas."
## [3] "Oakland, located on the South Side of Chicago, Illinois, USA, is one of 77 officially designated community areas in Chicago. It is located on the South Side of Chicago, Illinois, and is one of the city's 77 community areas."
## [4] "Fuller Park is the 37th of Chicago's 77 community areas. Located on the city's South Side, it is one of the city's 77 community areas."
## [5] "Kenwood, one of Chicago's 77 community areas, is on the shore of Lake Michigan on the South Side of Chicago. It is one of the city's 77 community areas."
## [6] "Washington Park is a community area on the South Side of Chicago which includes the 372 acre (1.5 square miles) area of the city's 77 community areas."
## [7] "Hyde Park is the 41st of the 77 community areas of Chicago. It is located on the South Side, near the city's 77 community areas."
## [8] "Grand Boulevard on the South Side of Chicago, Illinois, is one of the city's Community Areas. The neighborhood is located on the South Side of Chicago, Illinois, and is one of the city's 77 community areas."

```

```
head(links)
```

```

##           link
## 1 https://en.wikipedia.org/wiki/Armour_Square,_Chicago
## 2 https://en.wikipedia.org/wiki/Douglas,_Chicago
## 3 https://en.wikipedia.org/wiki/Oakland,_Chicago
## 4 https://en.wikipedia.org/wiki/Fuller_Park,_Chicago
## 5 https://en.wikipedia.org/wiki/Kenwood,_Chicago
## 6 https://en.wikipedia.org/wiki/Washington_Park,_Chicago_(neighborhood)
##      place      city
## 1  Armour Square    Chicago
## 2    Douglas      Chicago
## 3    Oakland      Chicago
## 4  Fuller Park     Chicago
## 5    Kenwood      Chicago
## 6 Washington Park Chicago (neighborhood)
##
## 1
## 2 Douglas, on the South Side of Chicago, Illinois, is one of Chicago's 77 community areas. The neighborhood is located on the South Side of Chicago, Illinois, and is one of the city's 77 community areas.
## 3
## 4
## 5
## 6 Kenwood, one of Chicago's 77 community areas, is on the shore of Lake Michigan on the South Side of Chicago. It is one of the city's 77 community areas.

```

Using a similar loop as in the last section, grab the descriptions of the various communities areas. Make a tibble with two columns: the name of the location and the text describing the location.

```
texts <- as_tibble(data.frame(links$place, links$desc))
```

Let's clean the data using `tidytext`. If you have trouble with this section, see the example shown in <https://www.tidytextmining.com/tidytext.html>

```
library(tidytext)
library(tm)
stop_words <- stopwords("en")
```

Create tokens using `unnest_tokens`. Make sure the data is in one-token-per-row format. Remove any stop words within the data. What are the most common words used overall?

```
words <- texts %>%
  unnest_tokens(word, links.desc)

words <- words %>%
  filter(!(word %in% stop_words))

head(texts)
```

```
## # A tibble: 6 x 2
##   links.place links.desc
##   <chr>      <chr>
## 1 Armour Square "Armour Square is a Chicago neighborhood on the city's South ~
## 2 Douglas      "Douglas, on the South Side of Chicago, Illinois, is one of C~
## 3 Oakland      "Oakland, located on the South Side of Chicago, Illinois, USA~
## 4 Fuller Park  "Fuller Park is the 37th of Chicago's 77 community areas. Loc~
## 5 Kenwood      "Kenwood, one of Chicago's 77 community areas, is on the shor~
## 6 Washington Park "Washington Park is a community area on the South Side of Chi~
```

```
words %>%
  count(word, sort = TRUE)
```

```
## # A tibble: 230 x 2
##   word      n
##   <chr>   <int>
## 1 chicago    14
## 2 south      14
## 3 community   11
## 4 park        11
## 5 area        10
## 6 side        10
## 7 douglas     7
## 8 kenwood     7
## 9 areas       6
## 10 chicago's  6
## # i 220 more rows
```

- “chicago” and “south” are the most common words used overall.

Plot the most common words within each location. What are some of the similarities between the locations? What are some of the differences?

```
words %>%
  aggregate(word ~ links.place, FUN = max)
```

```
##      links.place  word
## 1  Armour Square  west
## 2      Douglas  whose
## 3  Fuller Park  within
## 4 Grand Boulevard  west
## 5      Hyde Park  south
## 6      Kenwood  years
## 7      Oakland   usa
## 8 Washington Park  west
```

- There are 5 places whose common words are related to direction.

```
library(ggplot2)

words %>%
  count(word, sort = TRUE) %>%
  filter(n > 2) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word)) +
  geom_col() +
  labs(y = NULL)
```

